

A collection of various geometric shapes and patterns, including circles, squares, rectangles, triangles, and polygons, some filled with solid colors (purple, blue, green) and others with patterns (dots, lines, wavy lines).

Problem Statement

- Your client is an Insurance company that has provided Health Insurance to its customers now they need your help in building a model to predict whether the policyholders (customers) from past year will also be interested in Vehicle Insurance provided by the company.
- An insurance policy is an arrangement by which a company undertakes to provide a guarantee of compensation for specified loss, damage, illness, or death in return for the payment of a specified premium. A premium is a sum of money that the customer needs to pay regularly to an insurance company for this guarantee.
- Building a model to predict whether a customer would be interested in Vehicle Insurance is extremely helpful for the company because it can then accordingly plan its communication strategy to reach out to those customers and optimise its business model and revenue.
- Now, in order to predict, whether the customer would be interested in Vehicle insurance, you have information about demographics (gender, age, region code type), Vehicles (Vehicle Age, Damage), Policy (Premium, sourcing channel) etc.

Scope of this project

- Using this model, will try to predict whether a Probability of Customer being interested in Vehicle Loan in the test set will be **1 : Customer is interested, 0 : Customer is not interested** after the evaluation process.


Data set provided

- train.csv - 381109 records with 12 columns
- test.csv - 127037 records with 11 columns

Observation:

	count	mean	std	min	25%	50%	75%	max
Age	381109.0	38.822584	15.511611	20.0	25.0	36.0	49.0	85.0
Driving_License	381109.0	0.997869	0.046110	0.0	1.0	1.0	1.0	1.0
Region_Code	381109.0	26.388807	13.229888	0.0	15.0	28.0	35.0	52.0
Previously_Insured	381109.0	0.458210	0.498251	0.0	0.0	0.0	1.0	1.0
Annual_Premium	381109.0	30564.389581	17213.155057	2630.0	24405.0	31669.0	39400.0	540165.0
Policy_Sales_Channel	381109.0	112.034295	54.203995	1.0	29.0	133.0	152.0	163.0
Vintage	381109.0	154.347397	83.671304	10.0	82.0	154.0	227.0	299.0
Response	381109.0	0.122563	0.327936	0.0	0.0	0.0	0.0	1.0

Correlation Heatmap



category and numeric features

Categorical Columns:

- ['Gender', 'Vehicle_Age', 'Vehicle_Damage']


Numerical Columns:

- ['Age', 'Driving_License', 'Region_Code', 'Previously_Insured', 'Annual_Premium', 'Policy_Sales_Channel', 'Vintage']

Model Prediction Using Base Model:

Using Logistic Regression:

Train Accuracy :	Train cls_rep				
0.8780269411290085		precision	recall	f1-score	support
Test Accuracy :	0	0.88	1.00	0.94	267700
0.8750623179659416	1	0.46	0.00	0.00	37187
Train AUC	accuracy				
0.500067599016647		macro avg	0.67	0.50	0.47
Valid AUC		weighted avg	0.83	0.88	0.82
0.5	Valid cls_rep				
Train cnf_matrix		precision	recall	f1-score	support
[[267693 7]	0	0.88	1.00	0.93	66699
[37181 6]]	1	0.00	0.00	0.00	9523
Valid cnf_matrix	accuracy				
[[66699 0]		macro avg	0.44	0.50	0.47
[9523 0]]		weighted avg	0.77	0.88	0.82



Train ROC_AUC Score 0.500067599016647


Test ROC_AUC Score 0.5

Model Comparison:

Using RandomForestClassifier:

```
Train Accuracy :  
0.7924870525801363  
Test Accuracy :  
0.7519876151242424  
Train AUC  
0.880140048881563  
Valid AUC  
0.7825851080938335  
Train cnf_matrix  
[[204578  63122]  
 [   146 37041]]  
Valid cnf_matrix  
[[49477 17222]  
 [ 1682  7841]]
```

	precision	recall	f1-score	support
0	1.00	0.76	0.87	267700
1	0.37	1.00	0.54	37187
accuracy			0.79	304887
macro avg	0.68	0.88	0.70	304887
weighted avg	0.92	0.79	0.83	304887
Valid cls rep				
	precision	recall	f1-score	support
0	0.97	0.74	0.84	66699
1	0.31	0.82	0.45	9523
accuracy			0.75	76222
macro avg	0.64	0.78	0.65	76222
weighted avg	0.89	0.75	0.79	76222



Hyperparameters Tuning for each model

Using StackingClassifier & GradientBoostingClassifier & XGBClassifier
& AdaboostClassifier:

Train Accuracy :
0.9986257203488506

Test Accuracy :
0.859541864553541

Train AUC
0.9966934479063045

Valid AUC
0.5569336412530882

Train cnf_matrix
[[267499 201]

[218 36969]]

Valid cnf_matrix
[[64054 2645]

[8061 1462]]

Train cls_rep


	precision	recall	f1-score	support
0	1.00	1.00	1.00	267700
1	0.99	0.99	0.99	37187

accuracy			1.00	304887
macro avg	1.00	1.00	1.00	304887
weighted avg	1.00	1.00	1.00	304887

Valid cls_rep

	precision	recall	f1-score	support
0	0.89	0.96	0.92	66699
1	0.36	0.15	0.21	9523

accuracy			0.86	76222
macro avg	0.62	0.56	0.57	76222
weighted avg	0.82	0.86	0.83	76222



Hyperparameters Tuning for each model


Using StackingClassifier & GradientBoostingClassifier & XGBClassifier
& AdaboostClassifier & LGBMClassifier:

```
[LightGBM] [Warning] Unknown parameter: eval_metric
[LightGBM] [Warning] Unknown parameter: eval_metric
Train Accuracy :
0.9986421198673607
Test Accuracy :
0.859397549263992
Train AUC
0.9966101643463174
Valid AUC
0.5530254928638304
Train cnf_matrix
[[267512    188]
 [    226 36961]]
Valid cnf_matrix
[[64128  2571]
 [ 8146  1377]]
Train cls_rep
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	267700
1	0.99	0.99	0.99	37187
accuracy			1.00	304887
macro avg	1.00	1.00	1.00	304887
weighted avg	1.00	1.00	1.00	304887

Valid cls rep

	precision	recall	f1-score	support
0	0.89	0.96	0.92	66699
1	0.35	0.14	0.20	9523
accuracy			0.86	76222
macro avg	0.62	0.55	0.56	76222
weighted avg	0.82	0.86	0.83	76222



Hyperparameters Tuning for each model


Using StackingClassifier & XGBClassifier & CatBoostClassifier:

```
Train Accuracy :  
0.7924542535431159  
Test Accuracy :  
0.7511217233869486  
Train AUC  
0.880086637867823  
Valid AUC  
0.7820003318552217  
Train cnf_matrix  
[[204571  63129]  
 [   149 37038]]  
Valid cnf_matrix  
[[49413 17286]  
 [ 1684  7839]]  
Train cls_rep
```

	precision	recall	f1-score	support
0	1.00	0.76	0.87	267700
1	0.37	1.00	0.54	37187
accuracy			0.79	304887
macro avg	0.68	0.88	0.70	304887
weighted avg	0.92	0.79	0.83	304887

```
Valid cls rep
```

	precision	recall	f1-score	support
0	0.97	0.74	0.84	66699
1	0.31	0.82	0.45	9523
accuracy			0.75	76222
macro avg	0.64	0.78	0.65	76222
weighted avg	0.89	0.75	0.79	76222



Hyperparameters Tuning for each model

Using StackingClassifier & XGBClassifier & RandomForestClassifier:

Train Accuracy :

0.7916703565583314

Test Accuracy :

0.7508068536642964

Train AUC

0.8798602207830089

Valid AUC

0.7819104353101337

Train cnf_matrix

[[204313 63387]

[130 37057]]

Valid cnf_matrix

[[49387 17312]


[1682 7841]]

Train cls_rep

	precision	recall	f1-score	support
0	1.00	0.76	0.87	267700
1	0.37	1.00	0.54	37187
accuracy			0.79	304887
macro avg	0.68	0.88	0.70	304887
weighted avg	0.92	0.79	0.83	304887

Valid cls rep

	precision	recall	f1-score	support
0	0.97	0.74	0.84	66699
1	0.31	0.82	0.45	9523
accuracy			0.75	76222
macro avg	0.64	0.78	0.65	76222
weighted avg	0.89	0.75	0.79	76222






Classification Analysis:

Applied the following models and scores obtained

Model	ROC_AUC score
LogisticRegression	0.500056483
DecisionTreeClassifier	0.782063676
StackingClassifier	0.78218885
XGBClassifier	0.782163676
CatBoostClassifier	0.782227953
RandomForestClassifier	0.782597639

Result uploaded in the analytics vidhya website and the recieved score

Final Evaluation Metric : ROC_AUC score

CODE FILE	SOLUTION FILE	PRIVATE SCORE	PUBLIC SCORE	
 Code File	 Solution File	0.7825976387	0.7776676998	

Web API Creation End-Point

Cross-sell Prediction App

Gender

Male

Vehicle_Age

1-2 Year

Vehicle_Damage

Yes

Age

40.00

Driving_License

1.00

Region_Code

28.00

Previously_Insured

0.00

Annual_Premium

33762.00

Policy_Sales_Channel

7.00

Vintage

111.00

Predict

The predicted value is:

value

1

FastAPI Creation End-Point

The screenshot shows a web browser window with a REST client interface. The address bar displays "0.0.0.0/docs#/default/predict_predict_post". The interface includes a text area for the request body, an "Execute" button, and a "Responses" section. The request body is a JSON object with the following fields: "Gender": "Male", "Age": 40, "Driving_License": 1, "Region_Code": 28, "Previously_Insured": 0, "Vehicle_Age": "1-2 Year", "Vehicle_Damage": "Yes", "Annual_Premium": 33762, "Policy_Sales_Channel": 7, and "Vintage": 111. The "Responses" section shows a "Curl" command and a "Request URL" of "http://0.0.0.0/predict". The "Server response" section shows a "Code" of 200 and a "Response body" of {"Response": 1}. There is also a "Download" button for the response body.

Not Secure 0.0.0.0/docs#/default/predict_predict_post

Gmail YouTube Google Translate Imported

```
{
  "Gender": "Male",
  "Age": 40,
  "Driving_License": 1,
  "Region_Code": 28,
  "Previously_Insured": 0,
  "Vehicle_Age": "1-2 Year",
  "Vehicle_Damage": "Yes",
  "Annual_Premium": 33762,
  "Policy_Sales_Channel": 7,
  "Vintage": 111
}
```

Execute Clear

Responses

Curl

```
curl -X 'POST' \
  'http://0.0.0.0/predict' \
  -H 'accept: application/json' \
  -H 'Content-Type: application/json' \
  -d '{
    "Gender": "Male",
    "Age": 40,
    "Driving_License": 1,
    "Region_Code": 28,
    "Previously_Insured": 0,
    "Vehicle_Age": "1-2 Year",
    "Vehicle_Damage": "Yes",
    "Annual_Premium": 33762,
    "Policy_Sales_Channel": 7,
    "Vintage": 111
  }'
```

Request URL

http://0.0.0.0/predict

Server response

Code	Details
200	<p>Response body</p> <pre>{ "Response": 1 }</pre> <p>Download</p> <p>Response headers</p>

GCP CloudRun End-Points: FastApi and Streamlit

- <https://crosssellfastapi-222441656201.us-central1.run.app>
- <https://crosssellstreamlit-222441656201.us-central1.run.app>

Github Main URL:

<https://github.com/mohamedsuhaib88/MLHack2024/tree/main>

Github URL for FastAPI CI/CD using GCP CloudRun:

- <https://github.com/mohamedsuhaib88/crosssellstreamlit/tree/main>

Github URL for Streamlit CI/CD using GCP CloudRun:

- <https://github.com/mohamedsuhaib88/crosssellfastapi/tree/main>

DockerHub Image:

msuhaiba/fastapi-crosssell:1.1

msuhaiba/fastapi-crosssellweb:1.2

DockerHub URL:

<https://hub.docker.com/layers/msuhaiba/fastapi-crosssell/1.1/images/sha256-482dc6887664a9b2799d87cd34a20db28cd3a1c0c7d7feb8757e4cda2f2bd144?context=repo>